

Knowledge That **Walks**.

Get one constraint's hardest
knowledge out of a single head —
before it walks out the door.

The most fragile asset on your balance sheet
isn't in the financials. It's the judgment one
person holds and no system has ever
captured.

Name what you're capturing.

You've already named a constraint. Before you fix it — or hand any part of it to AI — you have to know whose head the knowledge lives in. **An agent is only as competent as the institutional context available to it.** This worksheet captures the knowledge one constraint depends on, before it walks.

1. Write your validated constraint in one line. Everything below scopes to it — not the whole company.
2. List the people whose **judgment** this constraint depends on. Run the walk test on each: *if they didn't show up tomorrow, what stops or degrades?*
3. Circle the one whose loss would hurt most. That's who you capture first.

THE CONSTRAINT — ONE SENTENCE (FROM SIGNAL)

e.g. Quotes take 5 days when they should take 1 — every RFQ routes through one person who reconciles pricing by hand.

THE WALK TEST — WHO HOLDS IT?

Whose judgment does this constraint depend on?

Who handles the exceptions · whose call can't be replaced by reading the SOP · who people walk over to and ask "how do I...?"

Person / role — what would stop or degrade if they were gone tomorrow

Person / role — what would stop or degrade if they were gone tomorrow

Person / role — what would stop or degrade if they were gone tomorrow

Person / role — what would stop or degrade if they were gone tomorrow

WORKED EXAMPLE — MERIDIAN (QUOTING CONSTRAINT)

Elena, VP Ops — the only person who can reconcile the rate card against 147 customer-specific pricing exceptions. If she's out, no quote goes out right.

Dave, Sr. Design Engineer (31 yrs) — estimates fabrication hours to within 10% by feel. Has never written the method down. Four years to retirement.

Score the **risk**.

Knowledge concentrated in one head is an **asset with a measurable cost** — and an exposure. Treat it like one. Score the person you circled in Step 1. The cost you can't put a number on is the cost that's invisible until it's gone.

1. Fill in every cell for your top person. Be specific — names, dates, dollars.
2. Flag **AT RISK** when concentration is **1** **and** an exit, retirement, role change, or reorg is anywhere on the horizon.
3. If you can't put a number on what it costs to lose them, you haven't found the real exposure yet.

DIMENSION	YOUR ANSWER
Who Name + role	<i>e.g. Elena Ruiz — VP Operations</i>
Concentration How many people could do this without them?	<i>e.g. Zero — she's the only one</i>
Exposure Planned exit, retirement, role change, reorg? When?	<i>e.g. No near-term exit, but no backup either</i>
Cost if it walks Hours to reconstruct · revenue at risk · errors introduced — a number	<i>e.g. ~2 years to rebuild + pricing errors across every quote</i>
Risk flag Concentration × exposure	<i>LOW / MEDIUM / AT RISK</i>

WORKED EXAMPLE — THE KNOWLEDGE THAT ALMOST WALKED

At a global food-safety company, a senior pricing strategist held **twenty years of customer-specific pricing exceptions** in his head — decisions made one at a time over decades, never consolidated into any system.

He wasn't flagged until the knowledge was scored. The team had **four weeks** before he retired. They captured everything. The loss they avoided would have taken **two years to reconstruct** and produced pricing errors throughout the transition. **AT RISK**

Sort what they **know**.

Not all knowledge is the same. Some can be written down and systematized — AI can run it. Some is judgment that has to stay human. Sort before you capture: it tells you what to turn into steps, what to turn into rules, and where a person has to stay in the loop. This is the **TML** framework — Task, Management, Leadership.

Task — documentable steps • AI runs it

Management — decision rights & rules • AI assists

Leadership — judgment & relationships • stays human

1. List what your top person actually knows about this constraint — one item per row.
2. Tag each as **T**, **M**, or **L**. Most "one person knows it" pain hides in M and L.
3. Tasks get written as steps. Leadership gets captured as *reasoning and examples* — never faked into a rule.

WHAT THEY KNOW	T / M / L
e.g. The 8-step process for reconciling a quote against the rate card	T
e.g. When a legacy customer gets exception pricing — and who must approve it	M
e.g. Reading whether a customer will walk over price	L

WHY THE SORT MATTERS

Point AI at the **Task** rows and it does the work directly. Give it the **Management** rules and it assists — flags the exception, routes the approval. The **Leadership** rows stay human; what you capture there is the *thinking*, so the next person can learn the judgment, not just the rule.

Capture it: the **interview**.

Don't ask them to write a summary — they'll compress and leave out the reasoning that makes the knowledge useful. **A transcript is the closest thing a person has to an API.** Sit them down, ask the real questions, and record it. The recording — transcribed — is the asset that feeds every workflow downstream.

1. Book **45–60 minutes** with your top person. Record it. Have it transcribed.
2. Work the questions below. Chase the *why* behind every answer — that's the part no SOP holds.
3. Store the transcript where both a person and an AI can reach it. A file on someone's desktop is not captured.

THE INTERVIEW SCRIPT

Ask these. Record the answers — don't summarize them.

- ☐ What decisions do you make that nobody else makes?
- ☐ Which exceptions have you handled that aren't written down anywhere?
- ☐ What rule of thumb do you use that you've never had to explain?
- ☐ Walk me through the last hard one. What did you weigh, and why?
- ☐ What would a new hire get wrong in their first month — and how would you catch it?

CAPTURE LOG	DETAIL
Who's interviewed	<i>Name + role</i>
Session booked for	<i>Date / time</i>
Recorded & transcribed?	<i>Yes / scheduled — tool used</i>
Where the transcript lives	<i>Shared, reachable location — not a desktop</i>

Before you call it **captured**.

Knowledge isn't captured because you held a meeting. It's captured when someone else — or an agent — can do the work from what you wrote down, without going back to the person. Run the check.

- ☐ **Someone else could do the work from this.**
No follow-up question to the original person required. If they'd still have to ask, it's not captured yet.
- ☐ **Task knowledge is written as steps** — not as prose, and not as "ask Elena."
A new hire could follow it on day one.
- ☐ **The judgment is captured as reasoning and examples** — not faked into a rule it isn't.
Leadership knowledge teaches the thinking, so the next person can learn it.
- ☐ **It lives where a person AND an AI can reach it.**
Not a head, not a desktop file, not a wiki nobody maintains.
- ☐ **Someone owns keeping it current.**
Name a knowledge manager — the person accountable for whether the knowledge base still reflects how the business runs today.

An agent is only as competent as the institutional context **available to it**.

What you just captured is the foundation every AI workflow on this constraint will stand on. Clean, current, reachable knowledge in — a well-briefed operator out. Skip it, and the AI runs confidently, plausibly, wrong — at scale. Knowledge that walks is a risk. Knowledge you captured is an asset.